What is Data Science?

Data Scientists: The Sexiest Job of the 21st Century

Source: Harvard Business Review

FiveThirtyEight

Data Science by FiveThirtyEight

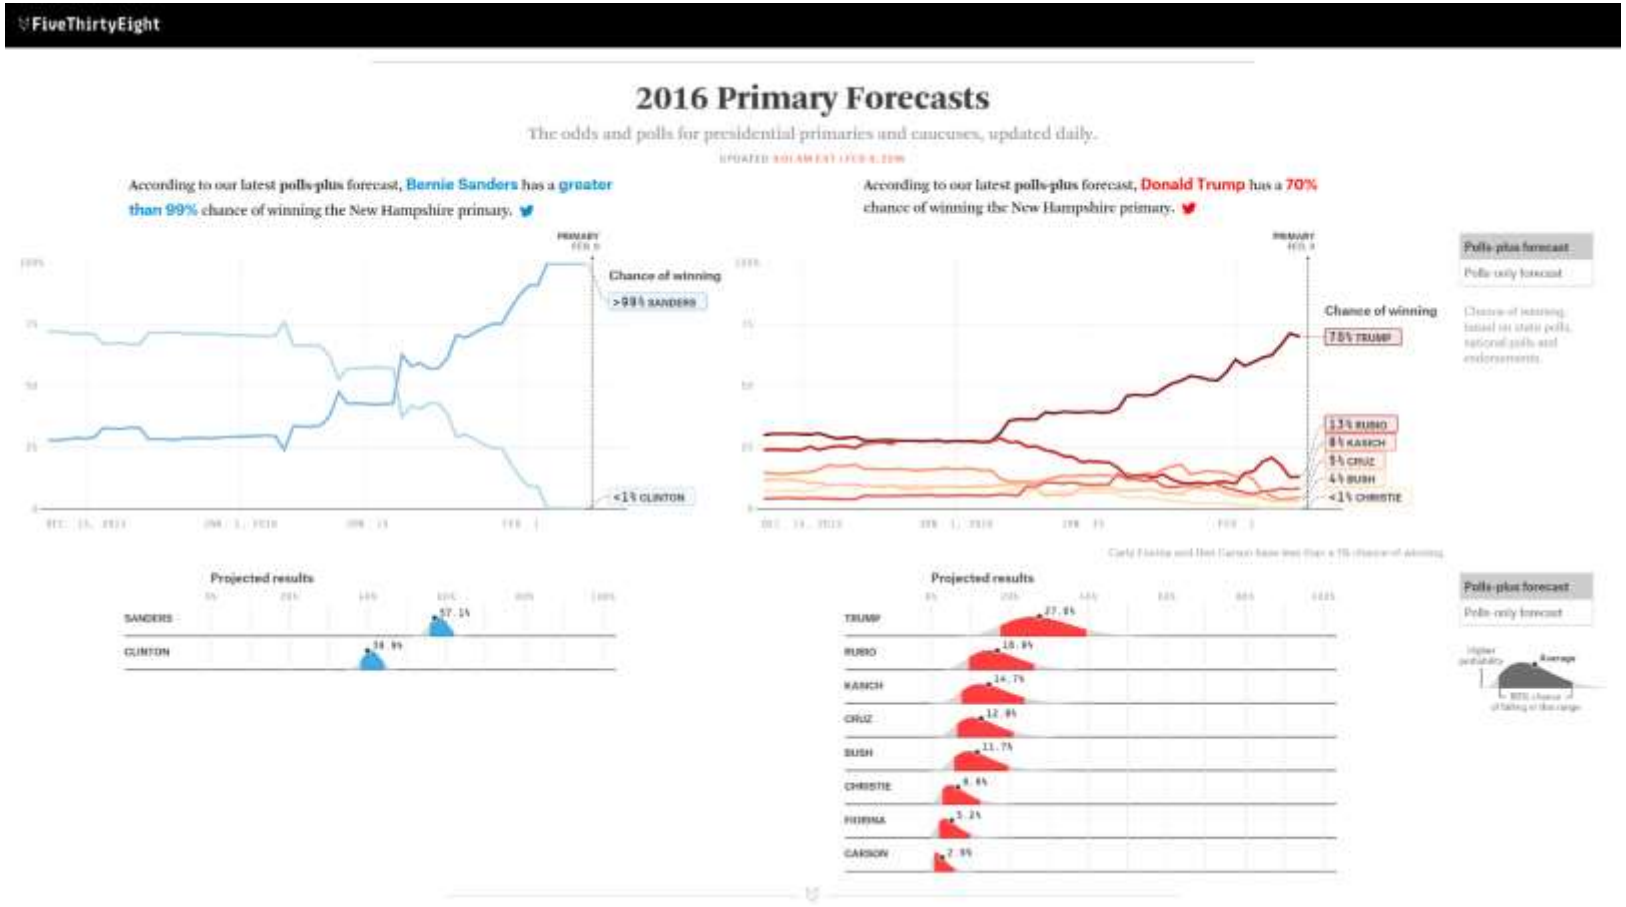

Source: FiveThirtyEight

Uber

Data Science by Uber

For Riders...

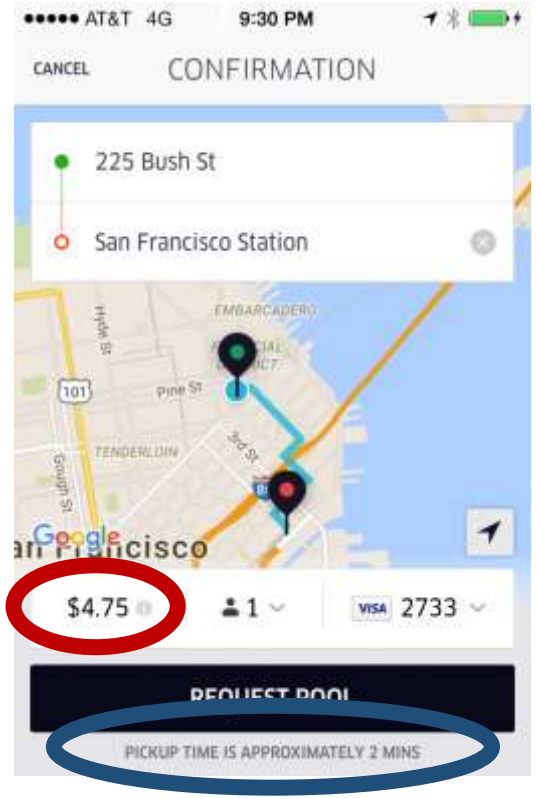

Fare Estimator and Dynamic Pricing (e.g., Surge Pricing) Algorithms

ETA (Estimated Time of Arrival) Algorithms

UBER

and Drivers!

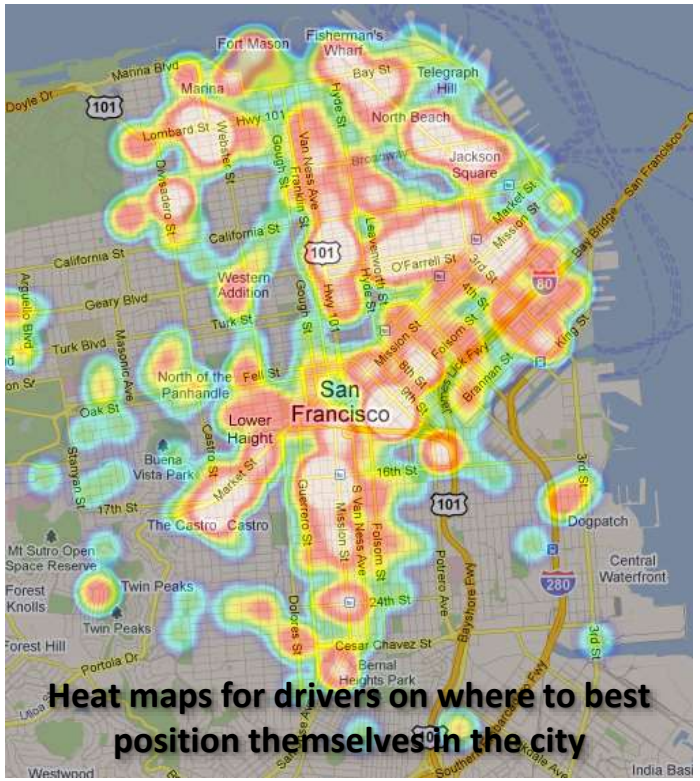

Sources: The Uber iPhone App/Wired

Data Science Based-Business Models is the New Normal

NETFLIX

amazon

UBER

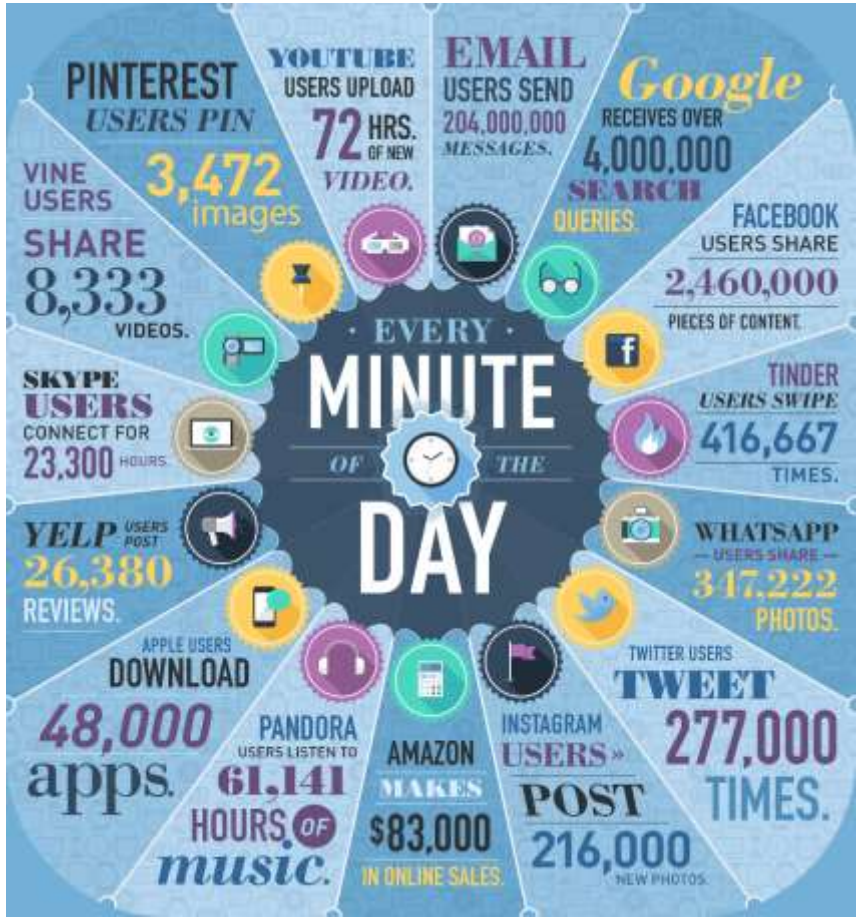
PINTEREST
USERS PIN
3,472 images
YOUTUBE
USERS UPLOAD
72 HRS. OF NEW VIDEO.
EMAIL
USERS SEND
204,000,000 MESSAGES.
Google
RECEIVES OVER
4,000,000
SEARCH QUERIES.
VINE USERS
SHARE
8,333
VIDEOS.
FACEBOOK
USERS SHARE
2,460,000
PIECES OF CONTENT.
EVERY
MINUTE
OF THE
DAY
SKYPE
USERS
CONNECT FOR
23,300 HOURS.
TINDER
USERS SWIPE
416,667
TIMES.
YELP USERS POST
26,380
REVIEWS.
WHATSAPP
USERS SHARE
347,222
PHOTOS.
APPLE USERS
DOWNLOAD
48,000
apps.
TWITTER USERS
TWEET
277,000
TIMES.
PANDORA
USERS LISTEN TO
61,141
HOURS OF
music.
AMAZON
MAKES
$83,000
IN ONLINE SALES.
INSTAGRAM
USERS
POST
216,000
NEW PHOTOS.


DOMO
DATA NEVER SLEEPS 2.0
How Much Data is Generated Every Minute?
THE GLOBAL INTERNET POPULATION GREW
14.3% FROM 2011-2013 AND NOW REPRESENTS
2.4 BILLION PEOPLE.
With each click, share and like, the world's data pool is expanding faster than we can comprehend. Businesses today are paying attention to scores of data sources to make crucial decisions about the future. The team at Domo can help your business make sense of this endless stream of data by providing executives with all their critical information in one intuitive platform. Domo delivers the insights you need to transform the way you run your business. Learn more at www.domo.com.
DOMO
SOURCES:

Source: DOMO

Who are Data Scientists?

Activity: Who are Data Scientists?

EXERCISE

DIRECTIONS (10 minutes)

1. Who are Data Scientists?
2. How do Data Scientists add value?
3. What makes a good Data Scientist?
4. When finished, share your answers with your table

DELIVERABLE

Answers to the above questions

Data Scientists in ≤140 characters

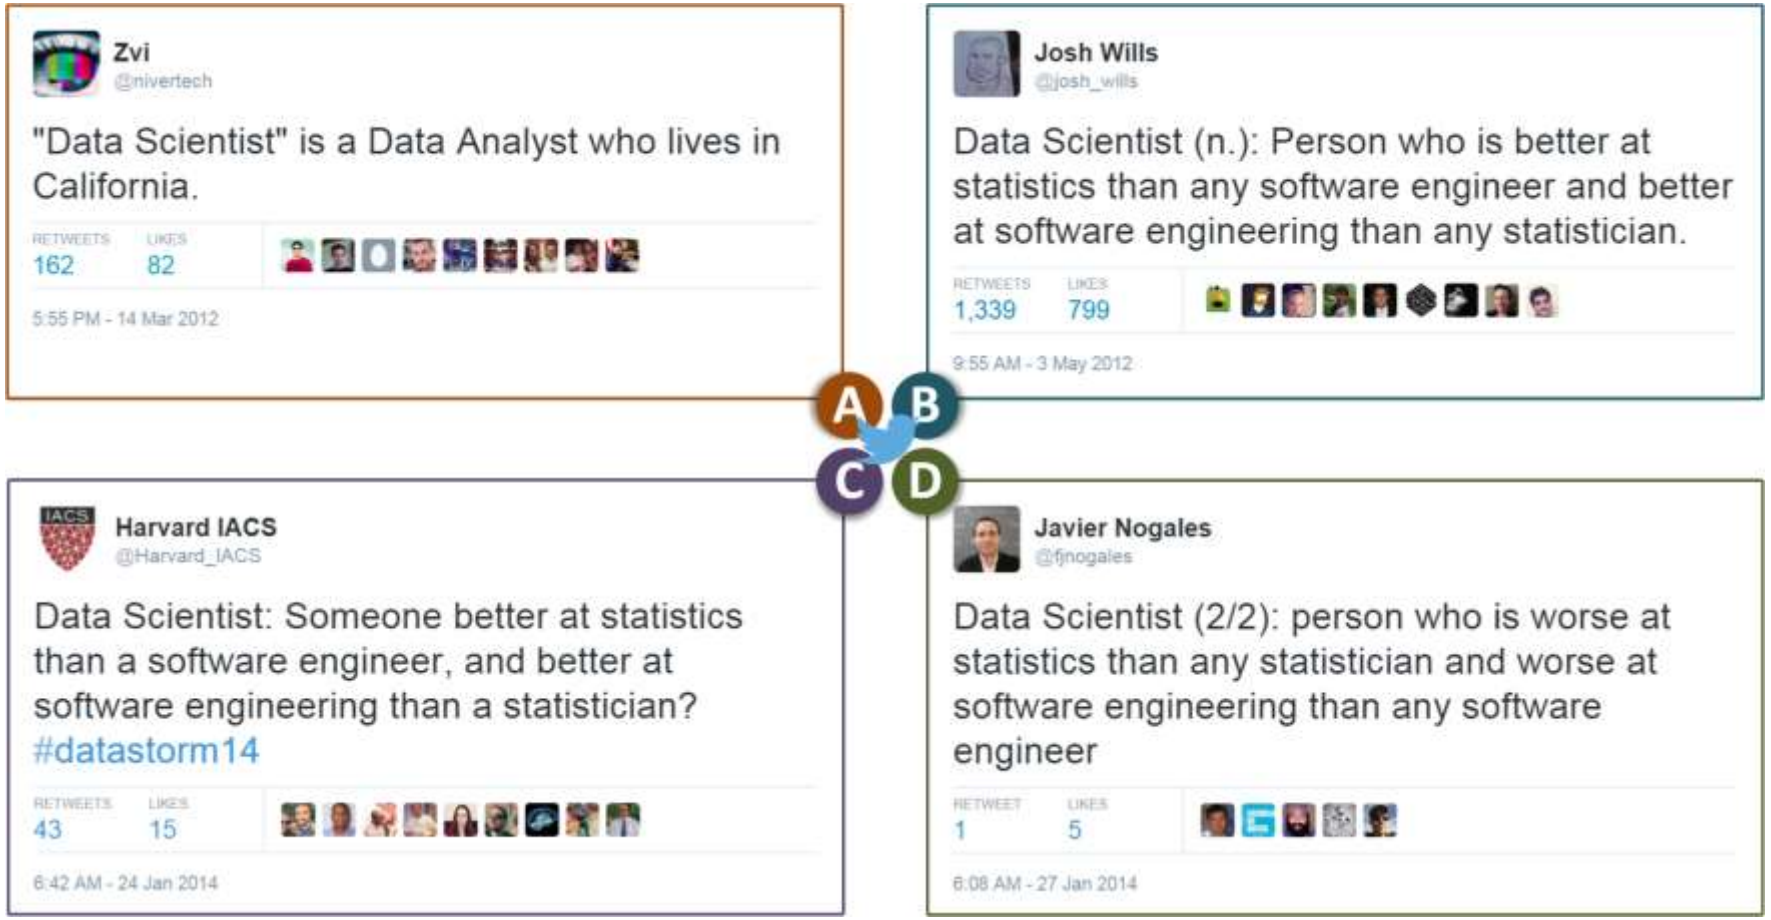

Source: Twitter

What is Data Science? (cont.)

Data is the New Oil of the Digital Economy and IoT, Big Data, DS, and Cloud relate to one another

Kirill Kirsanov/Shao-Chun Wang © 123RF.com

Data Science involves a variety of skillsets, not just one

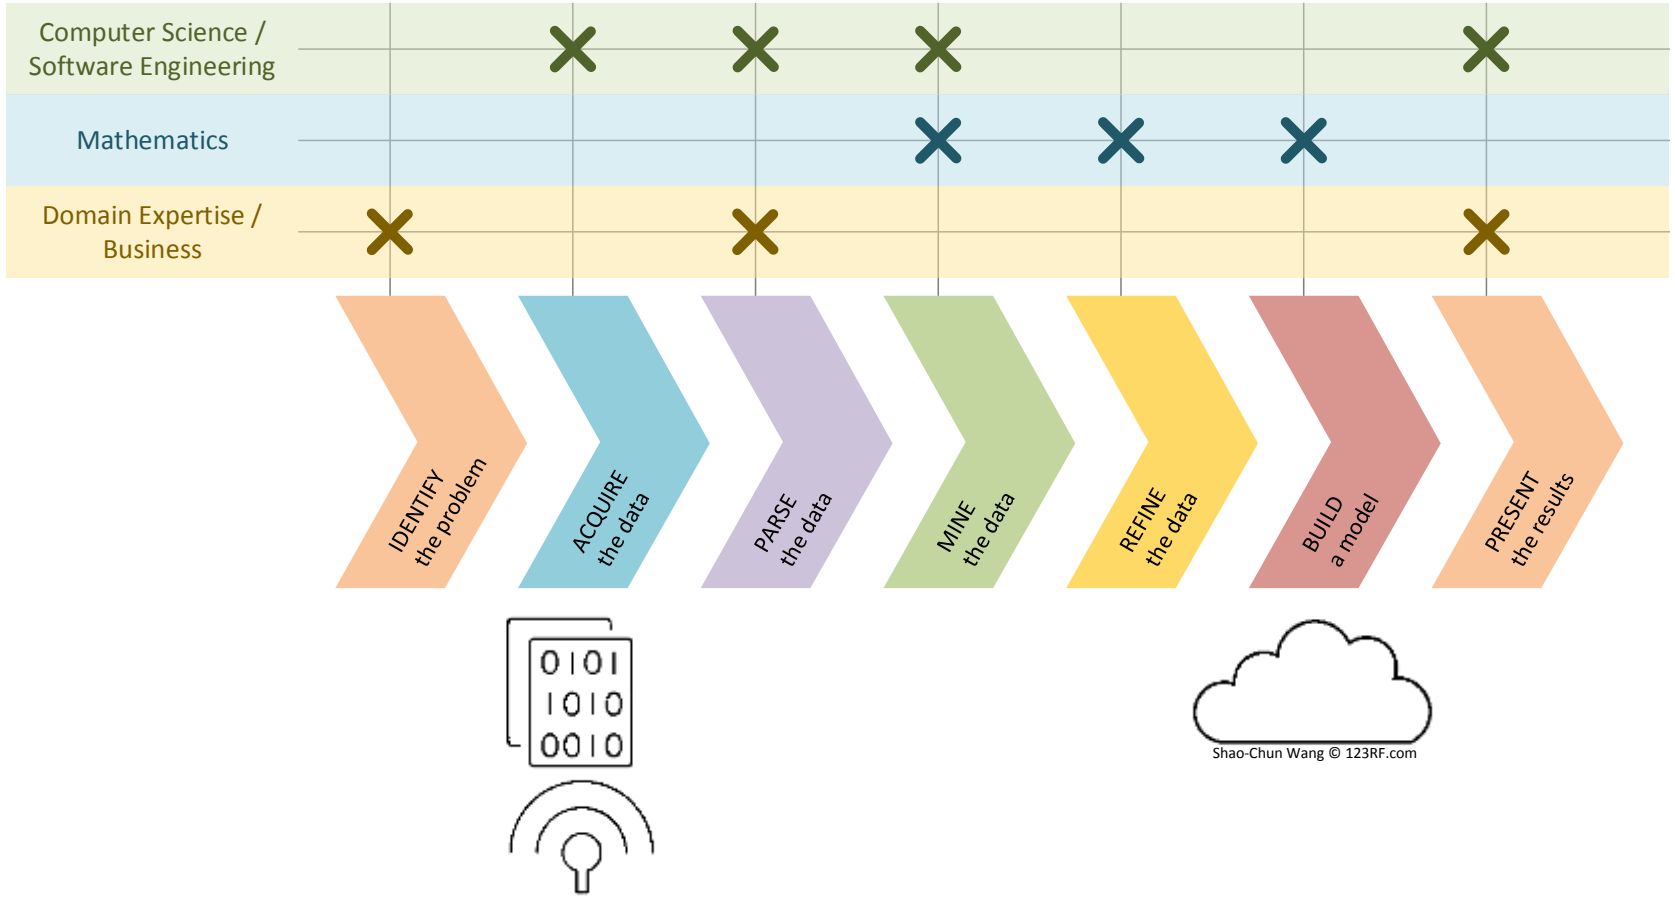

Data Science involves a variety of skillsets, not just one (cont.)

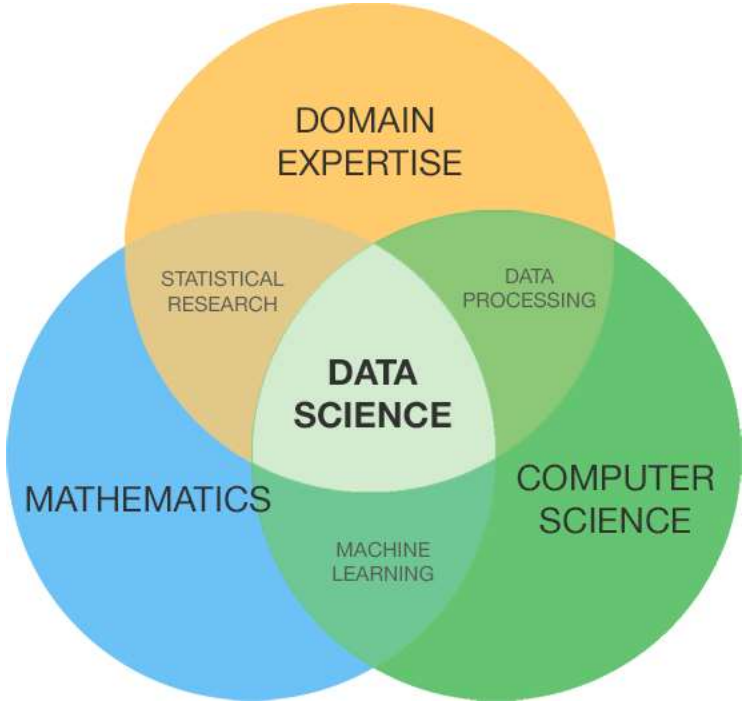

Source: Data Science for the C-suite

Data Scientists have different roles that prioritize different skillsets but all roles involve some part of each skillset to form strong data teams

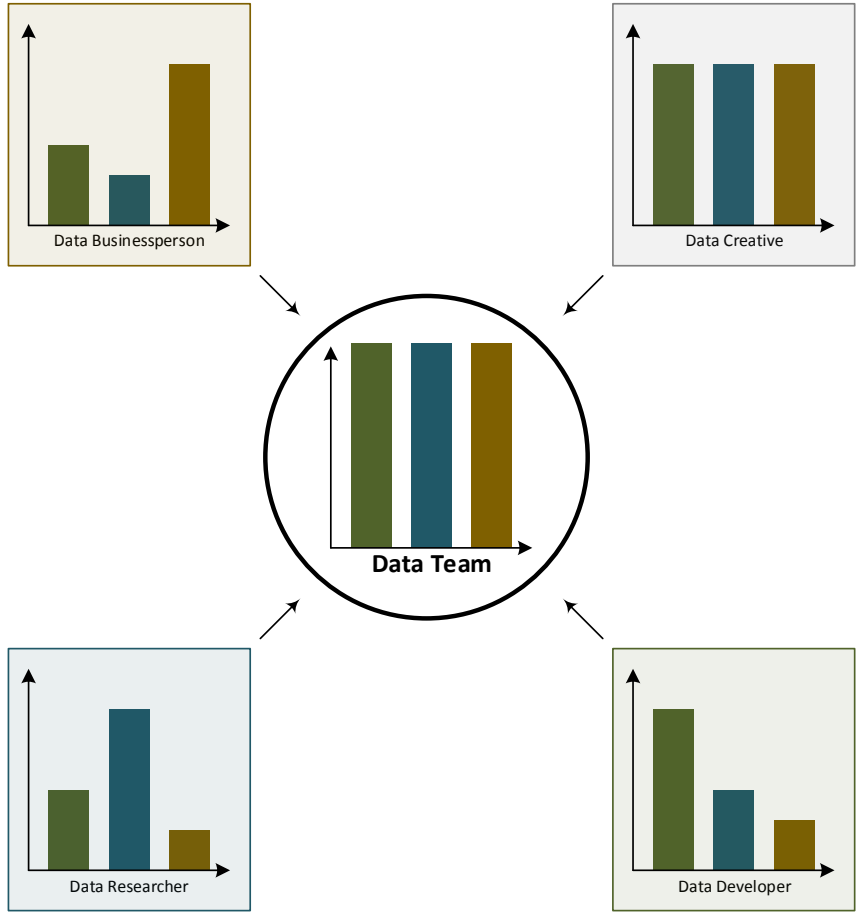

To sum it up

- Data Science is an interdisciplinary field about processes and systems to extract knowledge or insights from data in various forms

- An (ideal) data scientist is "someone who has the both the engineering skills to acquire and manage large data sets, and also has the statistician's skills to extract value from the large data sets and present that data to a large audience" – John Rauser

Data Science Workflow (and how it maps to the course)

What is the Data Science Workflow for?

- A methodology for Data Science to produce *reliable* and *reproducible* results
  - **Reliable**: Accurate findings
  - **Reproducible**: Others can follow your steps and get the same results
- Similar to the scientific method

The scientific method:

- Ask a Question
- Do Background Research
- Construct a Hypothesis
- Test Your Hypothesis by Doing an Experiment
- Analyze Your Data and Draw a Conclusion
- Communicate Your Results

The Data Science Workflow (also called the Data Science Pipeline)

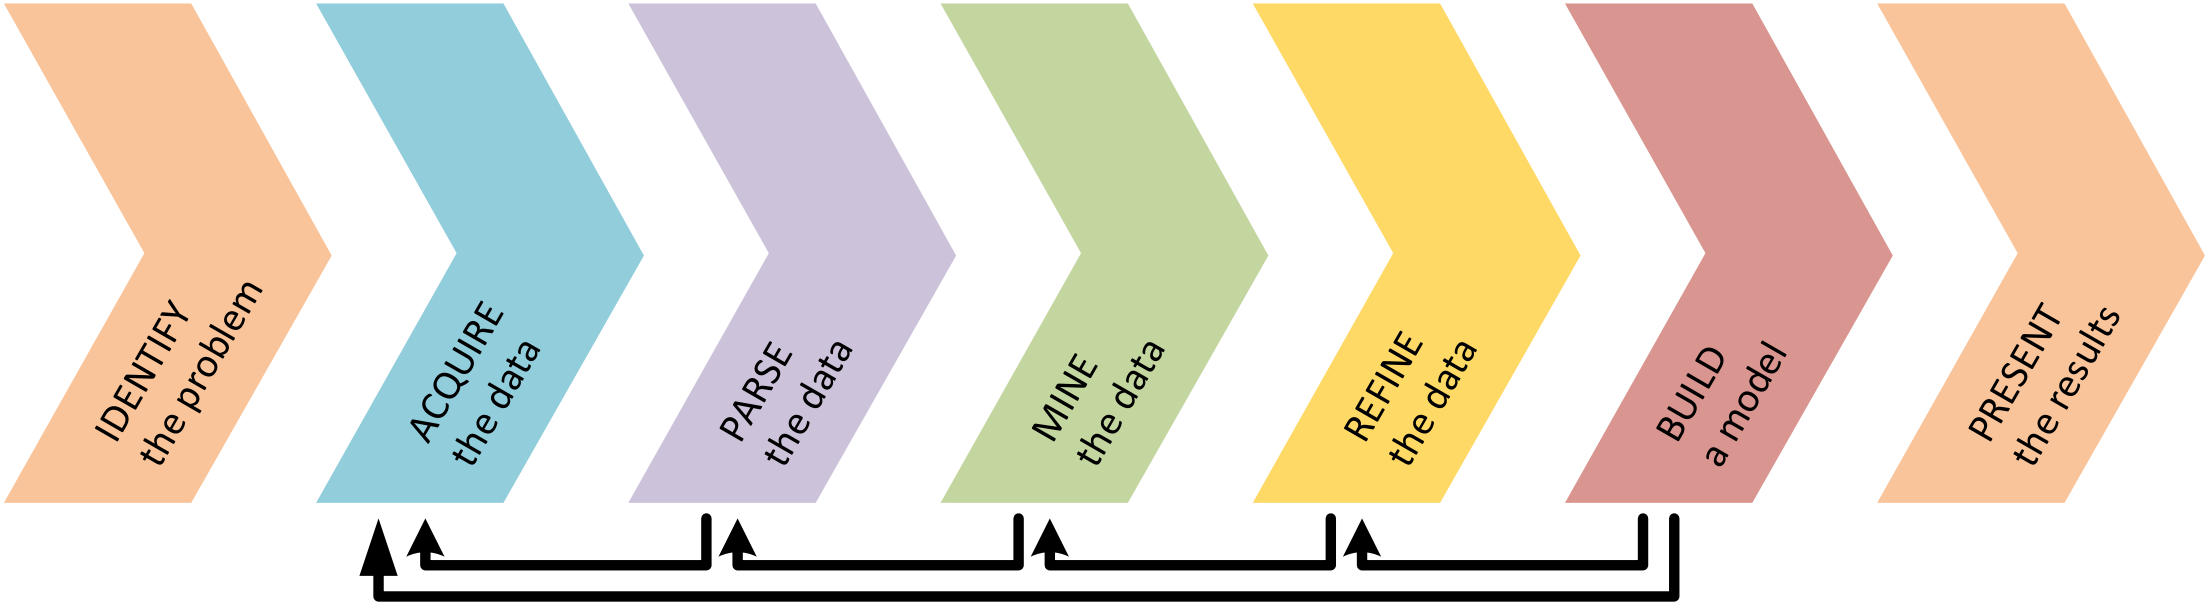

❶ Identify the Problem

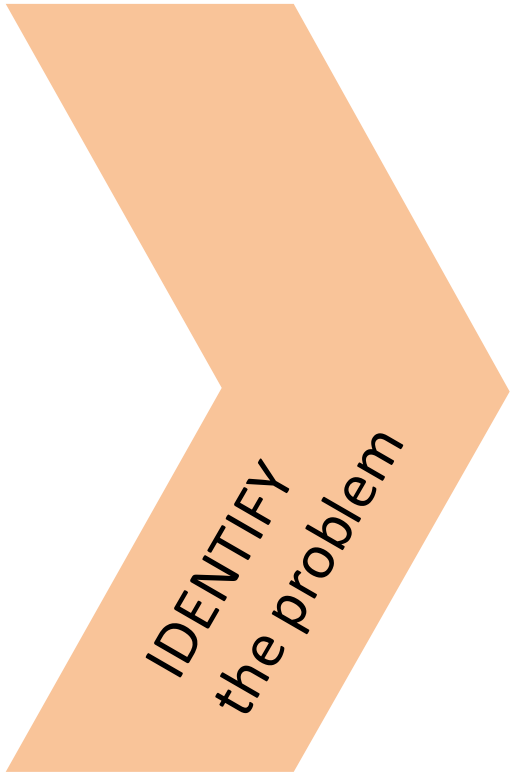

- Identify the Problem
  - Identify business/product objectives
  - Identify and hypothesize goals and criteria for success
  - Create a set of questions for identifying correct dataset

The Why’s and How’s of a Good Question

Corina Rosu © 123RF.com

❷ Acquire the Data

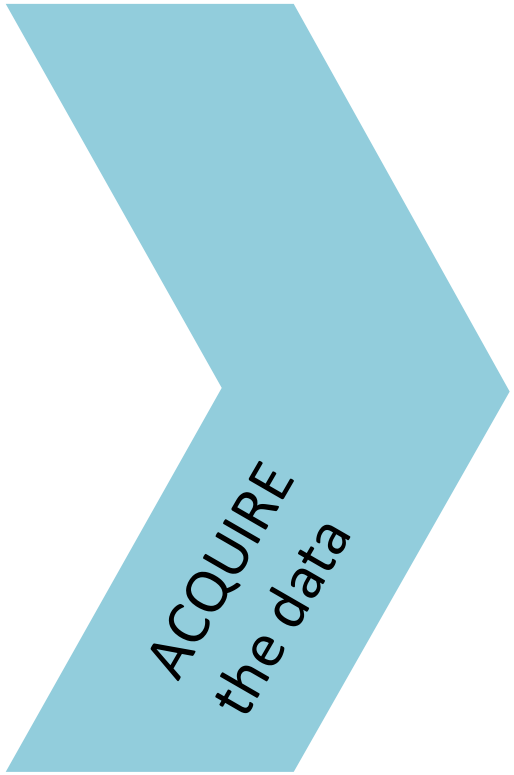

- Acquire the Data
  - Identify the “right” dataset(s)
  - Import data and set up local or remote data structure
  - Determine most appropriate tools to work with data

The data can be either unstructured or structured data

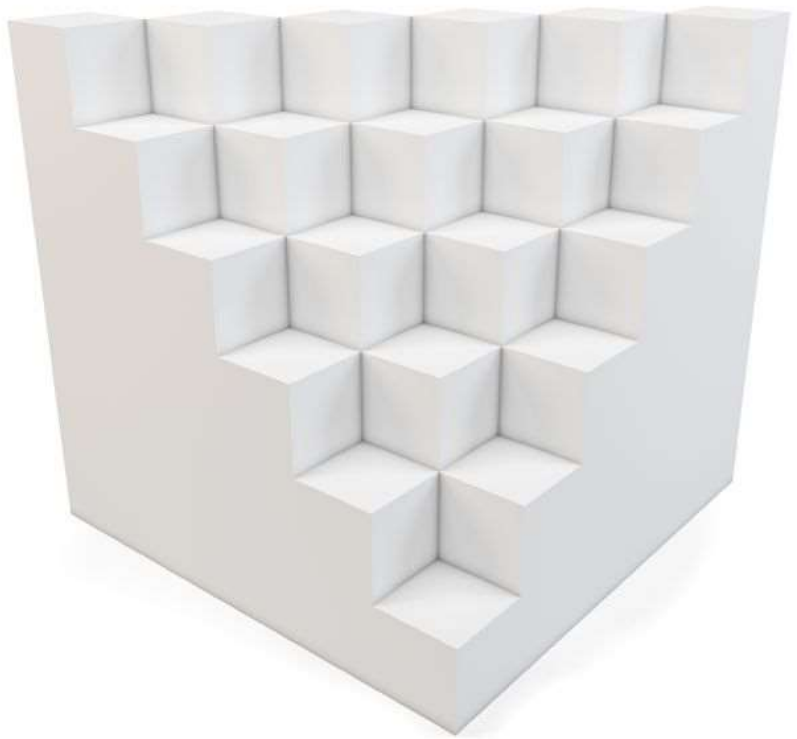

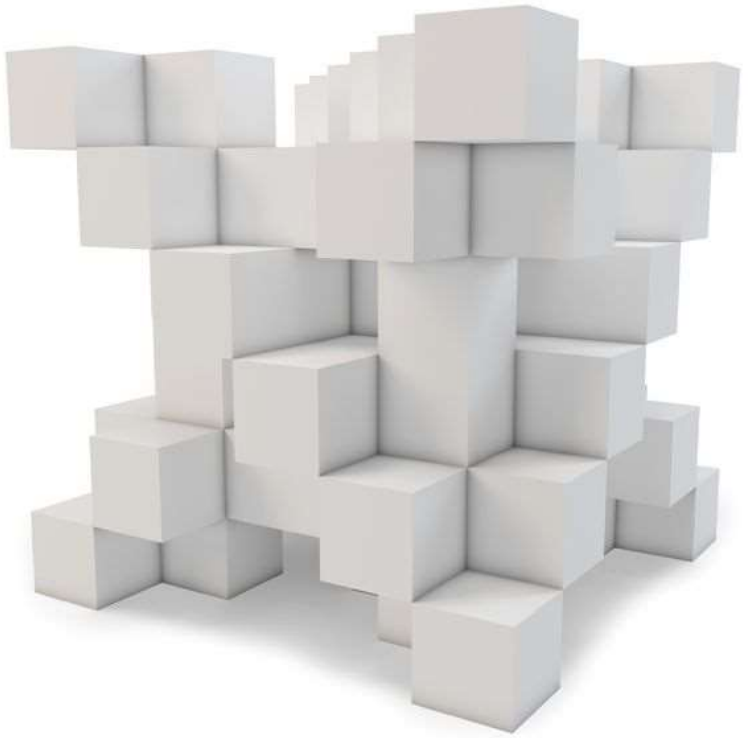

dny3d © 123RF.com

What's an example of unstructured data?

Bundit Chuangboonsri © 123RF.com

- Sessions 13 and 14 in Unit 3
  - Natural Language Processing

However, most of the course will focus on structured data

- Unit 2
  - Linear Regression (sessions 6 and 7)
  - Classification and Logistic Regression (session 8 and 9)
- Unit 3
  - Decision Trees and Random Forests (session 12)

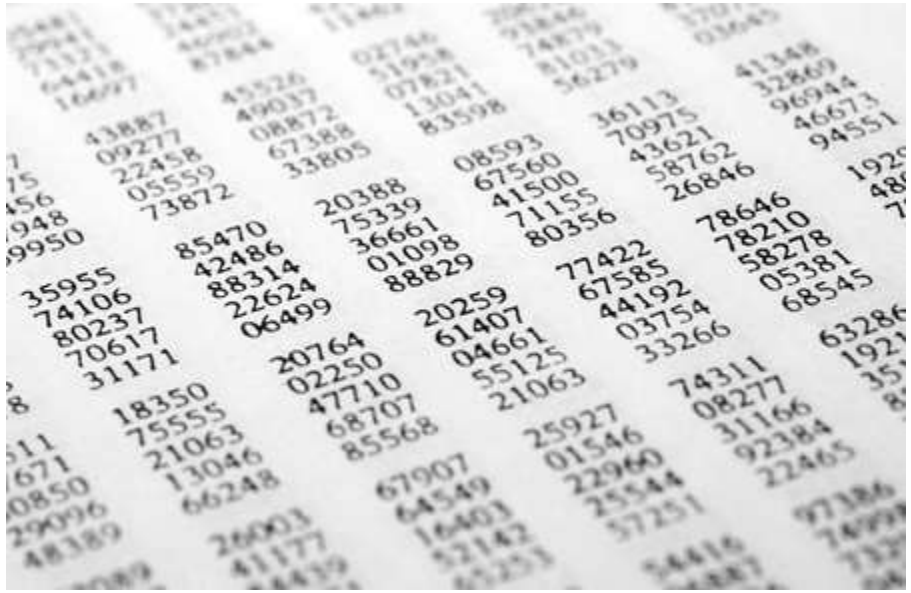

milosb © 123RF.com

Unstructured data can be longitudinal

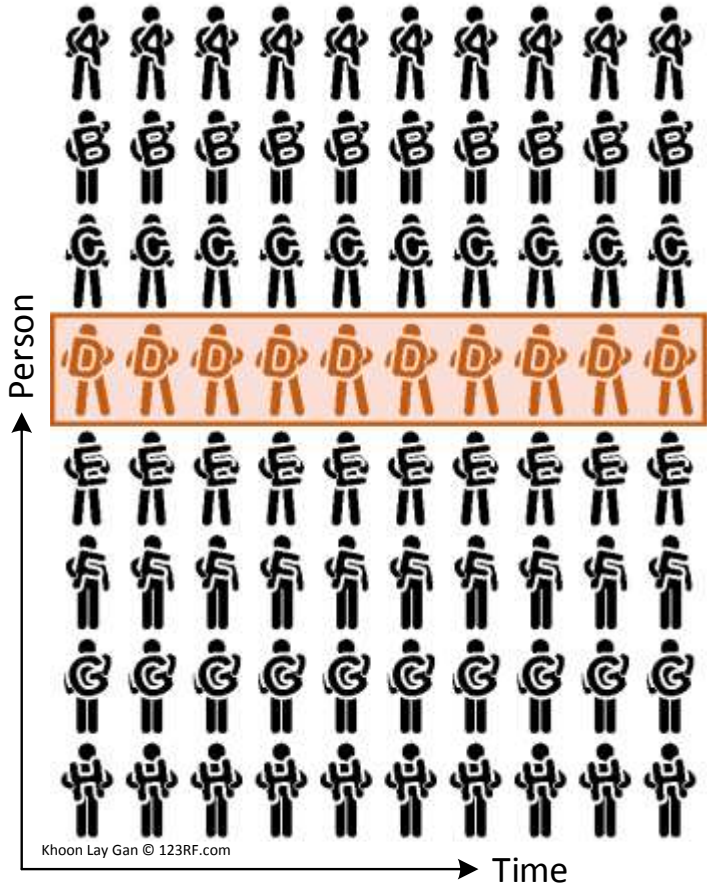

- Sessions 15 and 16 in Unit 3
  - Time Series

Unstructured data can be cross-sectional

- And most of the course will focus on it

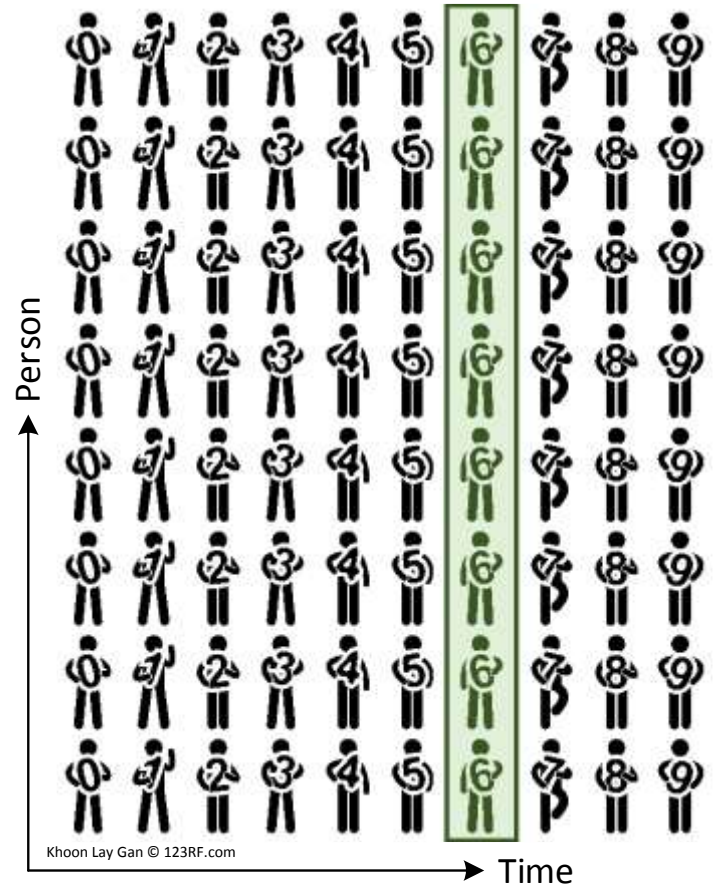

Raw structured data is Messy™…

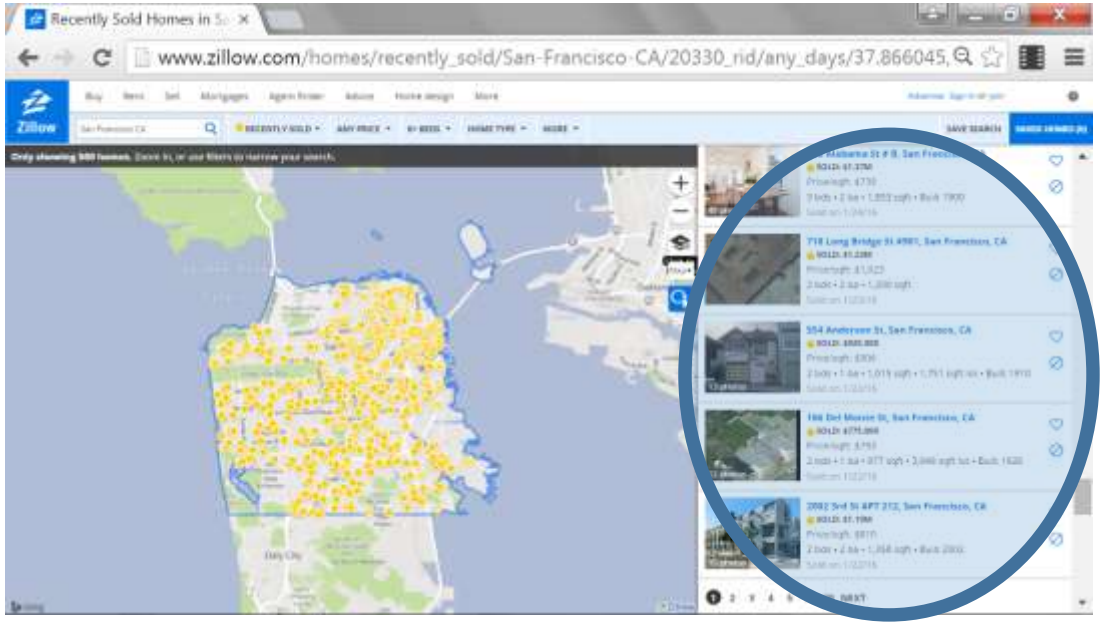


```
<div class="property-info"
id="yui_3_18_1_1_1456167242885_71870"><strong
id="yui_3_18_1_1_1456167242885_71869"><dt class="property-
address" id="yui_3_18_1_1_1456167242885_71868"><a
href="/homedetails/149-Shipley-St-San-Francisco-CA-
94107/15147894_zpid/" class="hdp-link routable" title="149
Shipley St, San Francisco, CA Real Estate"
id="yui_3_18_1_1_1456167242885_71873">149 Shipley St, San
Francisco, CA</a></dt></strong><dt class="listing-type zsg-
content_collapsed"
id="yui_3_18_1_1_1456167242885_71875"><span class="zsg-
icon-recently-sold type-icon"></span>Sold: $1.18M</dt><dt
class="zsg-fineprint"
id="yui_3_18_1_1_1456167242885_71877">Price/sqft:
$1,116</dt><dt class="property-data"
id="yui_3_18_1_1_1456167242885_71880"><span class="beds-
baths-sqft">3 bds • 2 ba • 1,057 sqft</span><span
class="built-year" id="yui_3_18_1_1_1456167242885_71879"> •
Built 1992</span></dt><dt class="sold-date zsg-fineprint"
id="yui_3_18_1_1_1456167242885_71975">Sold on
2/22/16</dt></div>
```

❸ Parse the Data

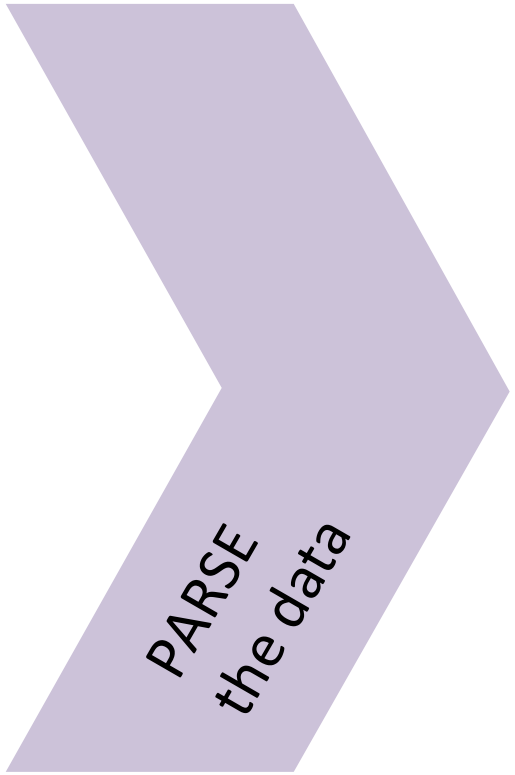

- Parse the Data
  - Read any documentation provided with the data
  - Perform exploratory data analysis
  - Verify the quality of the data

Exploratory Data Analysis

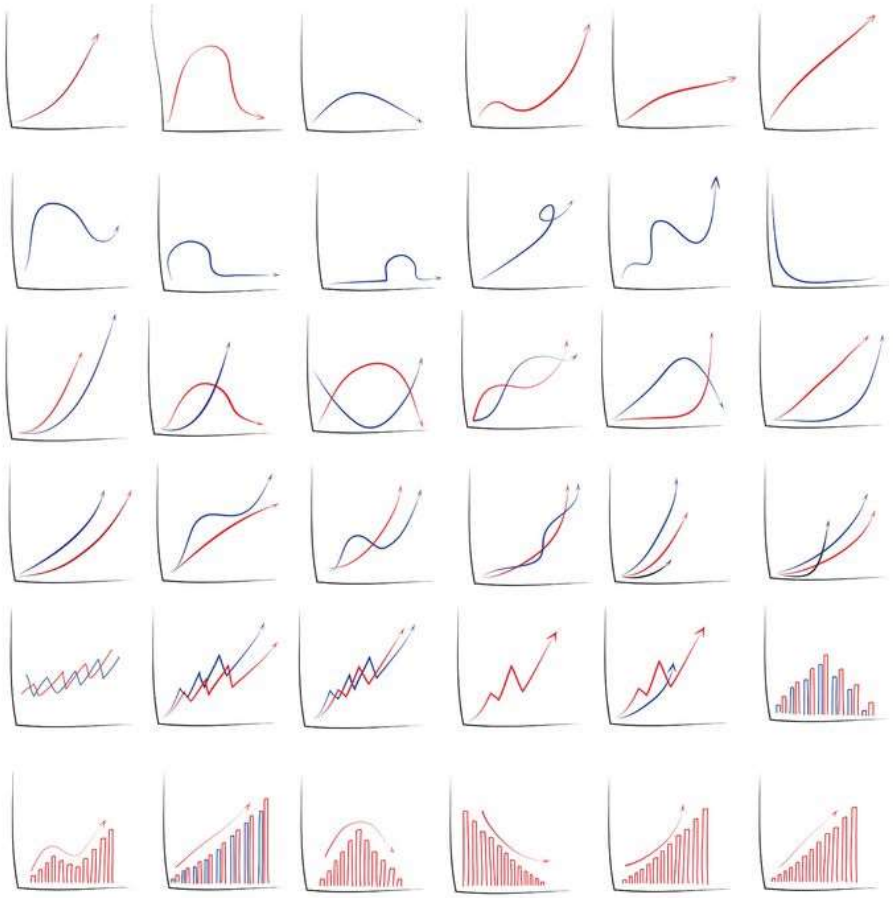

Napat Polchoke © 123RF.com

❹ Mine the Data

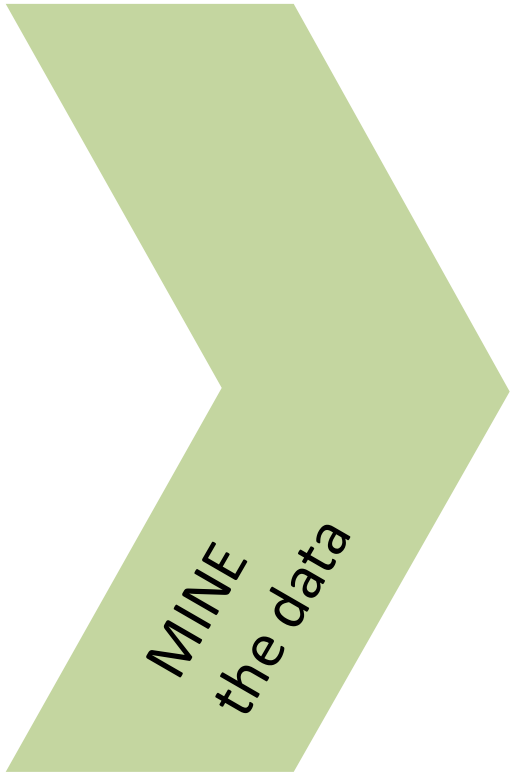

- Mine the Data
  - Determine sampling methodology and sample data
  - Format, clean, slice, and combine data in Python
  - Create necessary derived columns from the data (new data)

We will be tidying our data using the Python *pandas* library

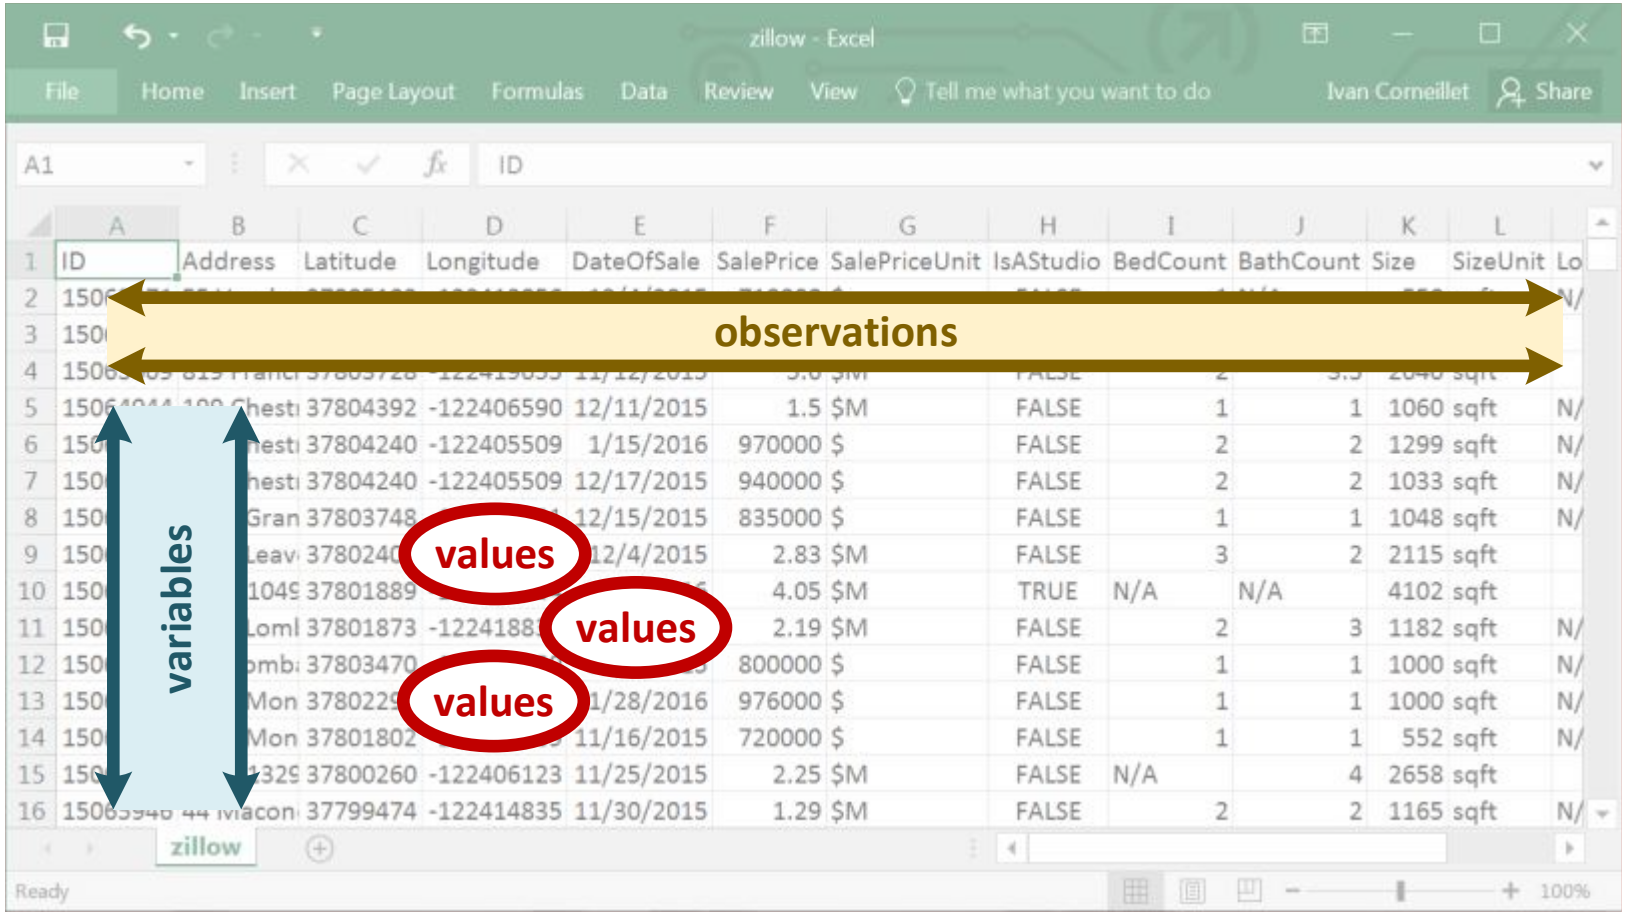

❺ Refine the Data

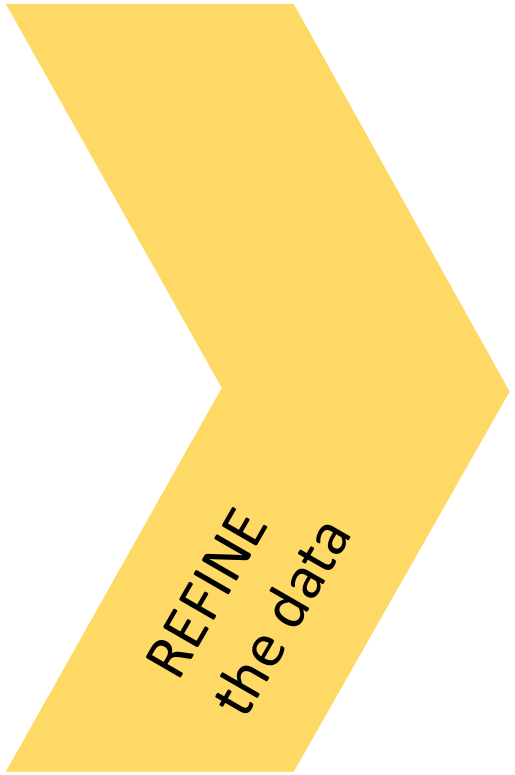

- Refine the Data
  - Identify trends and outliers
  - Apply descriptive and inferential statistics
  - Document and transform data

We will apply inferential statistics

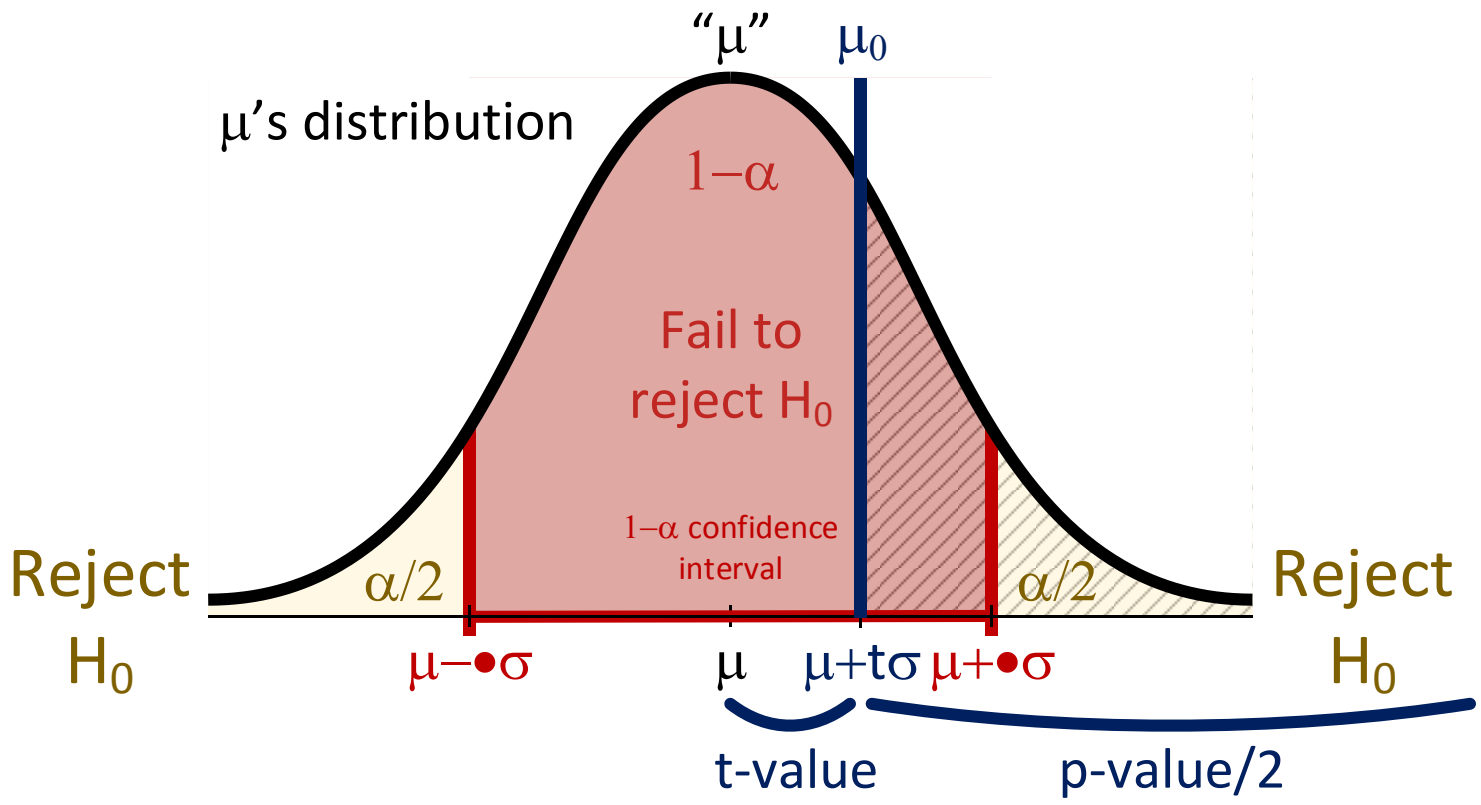

❻ Build a Model

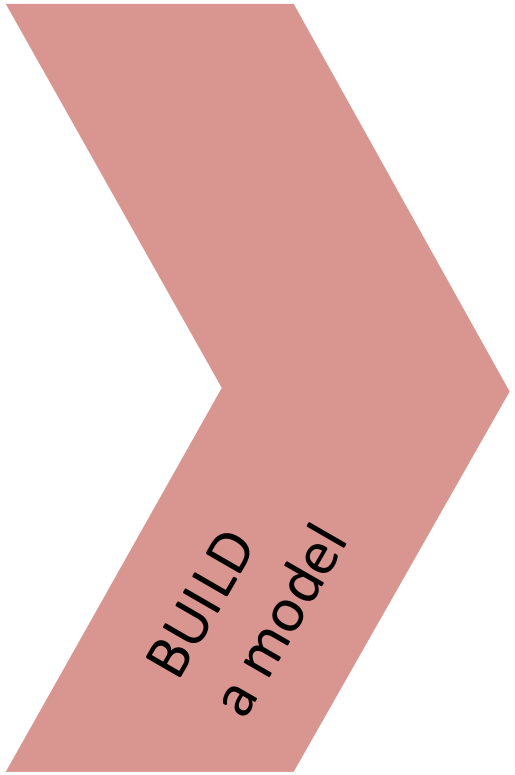

- Build a Model
  - Select appropriate model
  - Build model
  - Evaluate and refine model

The types of machine learning algorithms we will study in this course

	Continuous	Categorical
Supervised (a.k.a., predictive modeling)	Linear Regression K-Nearest Neighbors Decision Trees and Random Forests	Logistic Regression K-Nearest Neighbors Decision Trees and Random Forests
Unsupervised	*A machine learning model that doesn't use labeled data is called unsupervised. It extract structure from the data. Goal is "representation"*	

❼ Present the Results

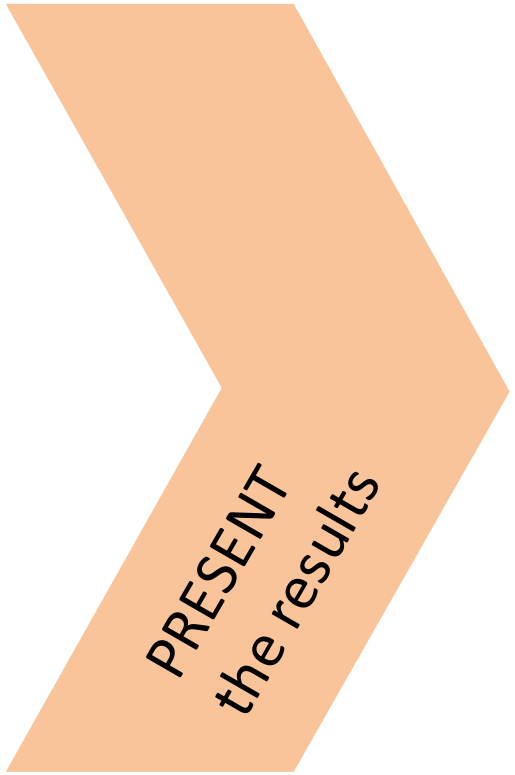

- Present the Results
  - Summarize findings with narrative, storytelling techniques
  - Present limitations and assumptions of your analysis
  - Identify follow up problems and questions for future analysis

Know Your Audience

A Note About Iteration

- Iteration is an important part of *every* step in the Data Science Workflow. At any given point in the process, you may find yourself repeating or going back and re-doing elements in order to better understand your data, clarify your model, and refine your presentation

- For example, after presenting your findings, you may want to:
  - Identify follow-up problems and questions for future analysis
  - Create a visually effective summary or report
  - Consider the needs of different stakeholders and how your report might be changed for them
  - Identify the limitations of your analysis
  - Identify relationships between visualizations

Multiple variants exist but they are pretty much all doing the same thing

- Jeff Hammerbacher
  - Identify problem
  - Instrument data sources
  - Collect data
  - Prepare data (integrate, transform, clean, impute, filter, aggregate)
  - Build model
  - Evaluate model
- Ben Fry
  - Acquire
  - Parse
  - Filter
  - Mine
  - Represent
  - Refine
  - Interact

- Peter Huber
  - Inspection
  - Error checking
  - Modification
  - Comparison
  - Modeling and model fitting
  - Simulation
  - What-if analyses
  - Interpretation
  - Presentation of conclusions
- Dataists
  - Obtain
  - Scrub
  - Explore
  - Model
  - Interpret

- Colin Mallows
  - Identify data to collect and its relevance to your problem
  - Statistical specification of the problem
  - Method selection
  - Analysis of method
  - Interpret results for non-statisticians
- Jim Gray
  - Capture
  - Curate
  - Communicate
- Ted Johnson
  - Assemble an accurate and relevant dataset
  - Choose the appropriate algorithm

Some great resources to follow along the class (or afterwards) *(optional; not required for the course)*

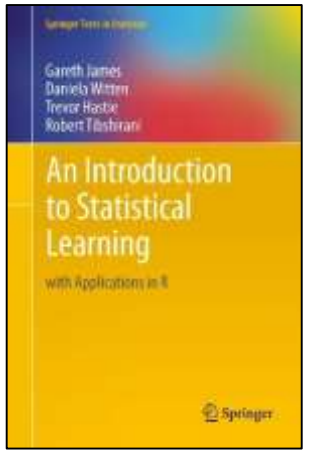

- An Introduction to Statistical Learning: with Applications in R (by James et al.). The e-book is available free-of-charge here

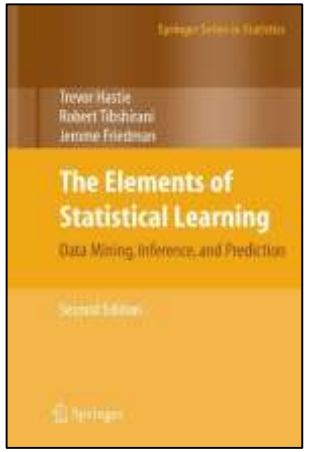

- For a more advanced treatment of these topics, check out The Elements of Statistical Learning: Data Mining, Inference, and Prediction (by Hastie et al.). And yes, the e-book is also free… (here)

Q & A

Onboarding Review

Python and iPython Notebooks

A couple of resources to get started with Python *(optional; not required for the course)* (the contents overlap but the styles are different so you should only pick one if any)

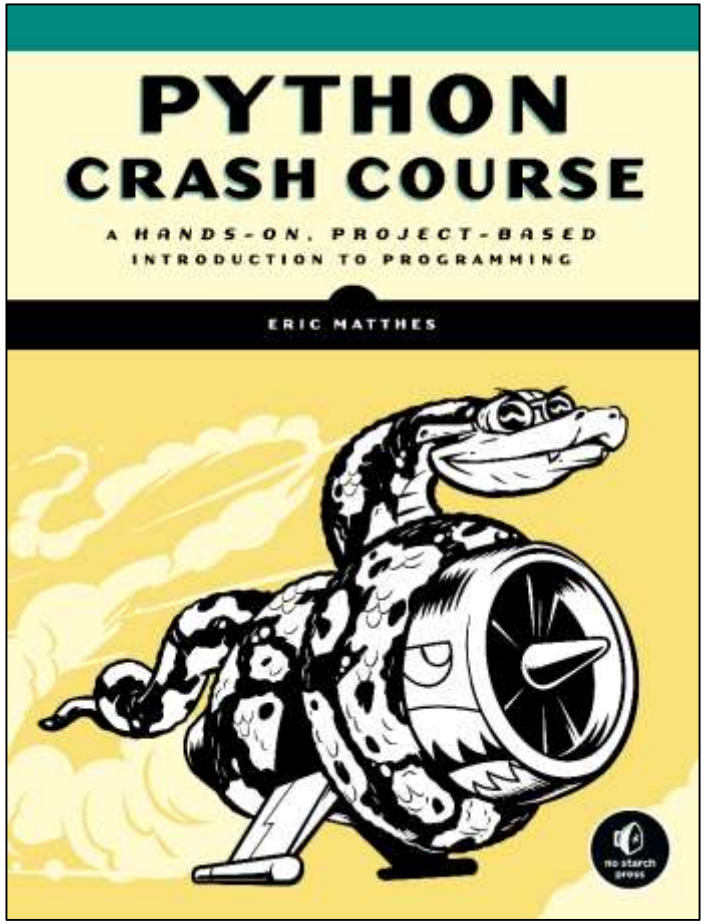

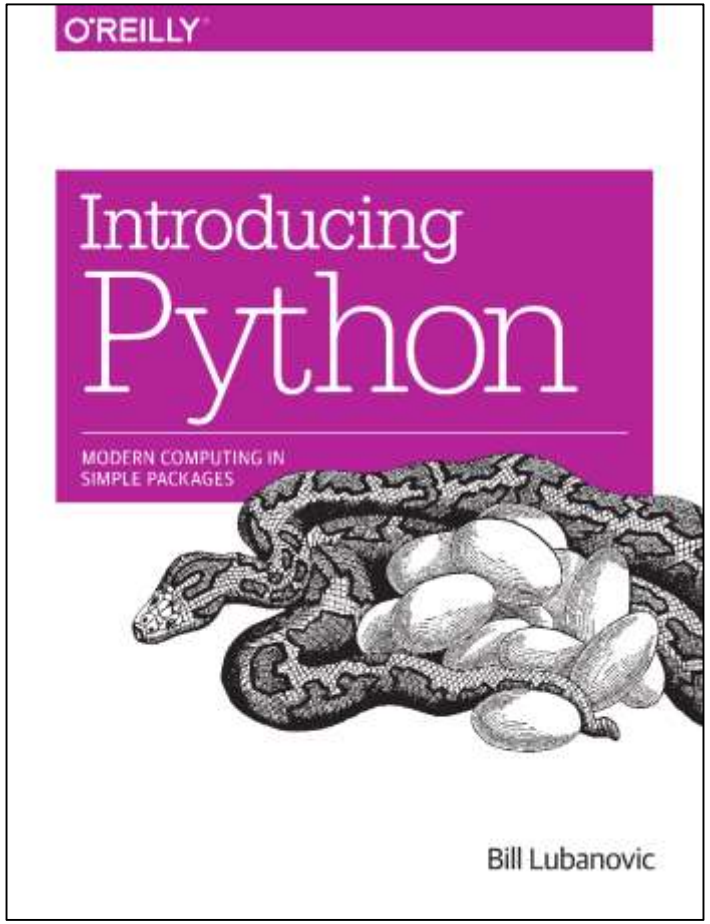

Review

Review

You should now be able to answer the following questions:

- What is data science?
- What is the data science workflow?
- How can you have a successful learning experience at GA?

Q & A

Next Class

Research Design and pandas

Learning Objectives

After the next lesson, you should be able to:

- Setup and manage your personal GitHub repository for submitting assignments
- Define a problem and types of data
- Identify dataset types
- Apply the data science workflow in the *pandas* context
- Write an iPython notebook to import, format, and clean data using the *pandas* library

Exit Ticket

Don’t forget to fill out your exit ticket here